

# Chitramoy\_Mukherjee-DSC630-Week-04-Assignment 2

April 7, 2024

DSC-630-T302

Chitramoy Mukherjee

Date : 04/04/2024

## Week3 - Exercise 4.2 - Clustering Exercise

```
[103]: from IPython.display import display, HTML
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.preprocessing import StandardScaler
from sklearn.cluster import KMeans
from sklearn.metrics import silhouette_score
from sklearn.decomposition import PCA
```

Import the dataset in pandas dataframe and Display the first few rows of the dataset.

```
[59]: df = pd.read_csv('C:
↳\\Users\\Chitramoy\\Desktop\\MS-DSC\\DSC-630\\Week-4\\als_data.csv')
print(df.info())
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2223 entries, 0 to 2222
Columns: 101 entries, ID to Urine.Ph_min
dtypes: float64(75), int64(26)
memory usage: 1.7 MB
None
```

```
[86]: # ID and SubjectID are the columns definitely not required for further analysis
↳process and can be considered as irrelevant column.
irrelevant_columns = ['ID', 'SubjectID']

# Remove irrelevant columns
data = df.drop(columns=irrelevant_columns)

# Describe data set post cleanup.
```

```
data.describe()
```

[86]:

	Age_mean	Albumin_max	Albumin_median	Albumin_min	Albumin_range \
count	2223.000000	2223.000000	2223.000000	2223.000000	2223.000000
mean	54.550157	47.011134	43.952542	40.766347	0.013779
std	11.396546	3.233980	2.654804	3.193087	0.009567
min	18.000000	37.000000	34.500000	24.000000	0.000000
25%	47.000000	45.000000	42.000000	39.000000	0.009042
50%	55.000000	47.000000	44.000000	41.000000	0.012111
75%	63.000000	49.000000	46.000000	43.000000	0.015873
max	81.000000	70.300000	51.100000	49.000000	0.243902

	ALSFRS_slope	ALSFRS_Total_max	ALSFRS_Total_median	ALSFRS_Total_min \
count	2223.000000	2223.000000	2223.000000	2223.000000
mean	-0.728274	31.692308	27.104926	19.877193
std	0.622329	5.314228	6.633643	8.583509
min	-4.345238	11.000000	2.500000	0.000000
25%	-1.086310	29.000000	23.000000	14.000000
50%	-0.620748	33.000000	28.000000	20.000000
75%	-0.283832	36.000000	32.000000	27.000000
max	1.207011	40.000000	40.000000	40.000000

	ALSFRS_Total_range	...	Sodium_median	Sodium_min	Sodium_range \
count	2223.000000	...	2223.000000	2223.000000	2223.000000
mean	0.026035	...	140.145254	136.755061	0.015000
std	0.016156	...	1.789886	2.715247	0.009283
min	0.000000	...	128.000000	112.000000	0.000000
25%	0.014035	...	139.000000	135.000000	0.010582
50%	0.023297	...	140.000000	137.000000	0.013123
75%	0.034799	...	141.000000	138.000000	0.017278
max	0.117647	...	146.500000	145.000000	0.142857

	trunk_max	trunk_median	trunk_min	trunk_range	Urine.Ph_max \
count	2223.000000	2223.000000	2223.000000	2223.000000	2223.000000
mean	6.203779	4.893387	2.955915	0.007136	6.820450
std	1.747660	2.146076	2.358095	0.004503	0.932141
min	0.000000	0.000000	0.000000	0.000000	5.000000
25%	5.000000	3.000000	1.000000	0.003643	6.000000
50%	7.000000	5.000000	3.000000	0.006920	7.000000
75%	8.000000	6.500000	5.000000	0.009639	7.000000
max	8.000000	8.000000	8.000000	0.042017	9.000000

	Urine.Ph_median	Urine.Ph_min
count	2223.000000	2223.000000
mean	5.710639	5.183221
std	0.625039	0.437222
min	5.000000	5.000000

25%	5.000000	5.000000
50%	6.000000	5.000000
75%	6.000000	5.000000
max	9.000000	8.000000

[8 rows x 99 columns]

```
[97]: import pandas as pd
from sklearn.preprocessing import StandardScaler

# Display basic information about the dataset
print(data.info())

# Apply StandardScaler to the numeric features
scaler = StandardScaler()
als_df_cleaned[numeric_features] = scaler.
    ↪fit_transform(als_df_cleaned[numeric_features])

als_df_cleaned.head()
```

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 2223 entries, 0 to 2222

Data columns (total 99 columns):

#	Column	Non-Null Count	Dtype
0	Age_mean	2223 non-null	int64
1	Albumin_max	2223 non-null	float64
2	Albumin_median	2223 non-null	float64
3	Albumin_min	2223 non-null	float64
4	Albumin_range	2223 non-null	float64
5	ALSFRS_slope	2223 non-null	float64
6	ALSFRS_Total_max	2223 non-null	int64
7	ALSFRS_Total_median	2223 non-null	float64
8	ALSFRS_Total_min	2223 non-null	int64
9	ALSFRS_Total_range	2223 non-null	float64
10	ALT.SGPT._max	2223 non-null	float64
11	ALT.SGPT._median	2223 non-null	float64
12	ALT.SGPT._min	2223 non-null	float64
13	ALT.SGPT._range	2223 non-null	float64
14	AST.SGOT._max	2223 non-null	int64
15	AST.SGOT._median	2223 non-null	float64
16	AST.SGOT._min	2223 non-null	float64
17	AST.SGOT._range	2223 non-null	float64
18	Bicarbonate_max	2223 non-null	float64
19	Bicarbonate_median	2223 non-null	float64
20	Bicarbonate_min	2223 non-null	float64
21	Bicarbonate_range	2223 non-null	float64
22	Blood.Urea.Nitrogen..BUN._max	2223 non-null	float64

23	Blood.Urea.Nitrogen..BUN._median	2223	non-null	float64
24	Blood.Urea.Nitrogen..BUN._min	2223	non-null	float64
25	Blood.Urea.Nitrogen..BUN._range	2223	non-null	float64
26	bp_diastolic_max	2223	non-null	int64
27	bp_diastolic_median	2223	non-null	float64
28	bp_diastolic_min	2223	non-null	int64
29	bp_diastolic_range	2223	non-null	float64
30	bp_systolic_max	2223	non-null	int64
31	bp_systolic_median	2223	non-null	float64
32	bp_systolic_min	2223	non-null	int64
33	bp_systolic_range	2223	non-null	float64
34	Calcium_max	2223	non-null	float64
35	Calcium_median	2223	non-null	float64
36	Calcium_min	2223	non-null	float64
37	Calcium_range	2223	non-null	float64
38	Chloride_max	2223	non-null	float64
39	Chloride_median	2223	non-null	float64
40	Chloride_min	2223	non-null	float64
41	Chloride_range	2223	non-null	float64
42	Creatinine_max	2223	non-null	float64
43	Creatinine_median	2223	non-null	float64
44	Creatinine_min	2223	non-null	float64
45	Creatinine_range	2223	non-null	float64
46	Gender_mean	2223	non-null	int64
47	Glucose_max	2223	non-null	float64
48	Glucose_median	2223	non-null	float64
49	Glucose_min	2223	non-null	float64
50	Glucose_range	2223	non-null	float64
51	hands_max	2223	non-null	int64
52	hands_median	2223	non-null	float64
53	hands_min	2223	non-null	int64
54	hands_range	2223	non-null	float64
55	Hematocrit_max	2223	non-null	float64
56	Hematocrit_median	2223	non-null	float64
57	Hematocrit_min	2223	non-null	float64
58	Hematocrit_range	2223	non-null	float64
59	Hemoglobin_max	2223	non-null	float64
60	Hemoglobin_median	2223	non-null	float64
61	Hemoglobin_min	2223	non-null	float64
62	Hemoglobin_range	2223	non-null	float64
63	leg_max	2223	non-null	int64
64	leg_median	2223	non-null	float64
65	leg_min	2223	non-null	int64
66	leg_range	2223	non-null	float64
67	mouth_max	2223	non-null	int64
68	mouth_median	2223	non-null	float64
69	mouth_min	2223	non-null	int64
70	mouth_range	2223	non-null	float64

71	onset_delta_mean	2223 non-null	int64
72	onset_site_mean	2223 non-null	int64
73	Platelets_max	2223 non-null	int64
74	Platelets_median	2223 non-null	float64
75	Platelets_min	2223 non-null	float64
76	Potassium_max	2223 non-null	float64
77	Potassium_median	2223 non-null	float64
78	Potassium_min	2223 non-null	float64
79	Potassium_range	2223 non-null	float64
80	pulse_max	2223 non-null	int64
81	pulse_median	2223 non-null	float64
82	pulse_min	2223 non-null	int64
83	pulse_range	2223 non-null	float64
84	respiratory_max	2223 non-null	int64
85	respiratory_median	2223 non-null	float64
86	respiratory_min	2223 non-null	int64
87	respiratory_range	2223 non-null	float64
88	Sodium_max	2223 non-null	float64
89	Sodium_median	2223 non-null	float64
90	Sodium_min	2223 non-null	float64
91	Sodium_range	2223 non-null	float64
92	trunk_max	2223 non-null	int64
93	trunk_median	2223 non-null	float64
94	trunk_min	2223 non-null	int64
95	trunk_range	2223 non-null	float64
96	Urine.Ph_max	2223 non-null	float64
97	Urine.Ph_median	2223 non-null	float64
98	Urine.Ph_min	2223 non-null	float64

dtypes: float64(75), int64(24)

memory usage: 1.7 MB

None

```
[97]:
```

	Age_mean	Albumin_max	Albumin_median	Albumin_min	Albumin_range	\
0	0.917137	3.089417	-1.300781	-0.866550	5.480929	
1	-0.574879	-0.622016	-1.112401	-0.553303	-0.347725	
2	-1.452535	0.924415	1.148162	1.326179	-0.507103	
3	0.741606	-0.003443	0.017880	0.073191	-0.174361	
4	0.741606	-0.003443	0.583021	0.386438	-0.573670	

	ALSFRS_slope	ALSFRS_Total_max	ALSFRS_Total_median	ALSFRS_Total_min	\
0	-0.381450	-0.318520	0.134960	0.247368	
1	-0.310907	0.998995	0.888863	0.130839	
2	-0.299769	-1.447819	-1.975969	-1.150976	
3	0.208801	-0.318520	0.285741	0.480425	
4	0.456831	0.057913	0.059570	0.014311	

	ALSFRS_Total_range	...	Sodium_median	Sodium_min	Sodium_range	\
--	--------------------	-----	---------------	------------	--------------	---

0	-0.301588	...	2.992342	2.300470	0.260968
1	0.166537	...	-1.198812	-0.278144	-0.489913
2	-0.064100	...	1.595291	1.195350	-0.654169
3	-0.685524	...	-0.639992	0.458603	-0.272701
4	-0.350529	...	-0.081171	0.458603	-0.722774

	trunk_max	trunk_median	trunk_min	trunk_range	Urine.Ph_max	\
0	1.028018	0.981832	1.715365	-0.997420	-0.880376	
1	1.028018	0.981832	0.867032	-0.388669	0.192665	
2	-0.688950	-2.280669	-1.253800	0.398249	-0.880376	
3	-0.688950	0.049689	0.018699	-0.477181	0.192665	
4	-0.116627	-0.416383	-0.829634	0.300598	-0.880376	

	Urine.Ph_median	Urine.Ph_min
0	0.463054	1.868532
1	-1.137208	-0.419151
2	-1.137208	-0.419151
3	0.463054	-0.419151
4	-1.137208	-0.419151

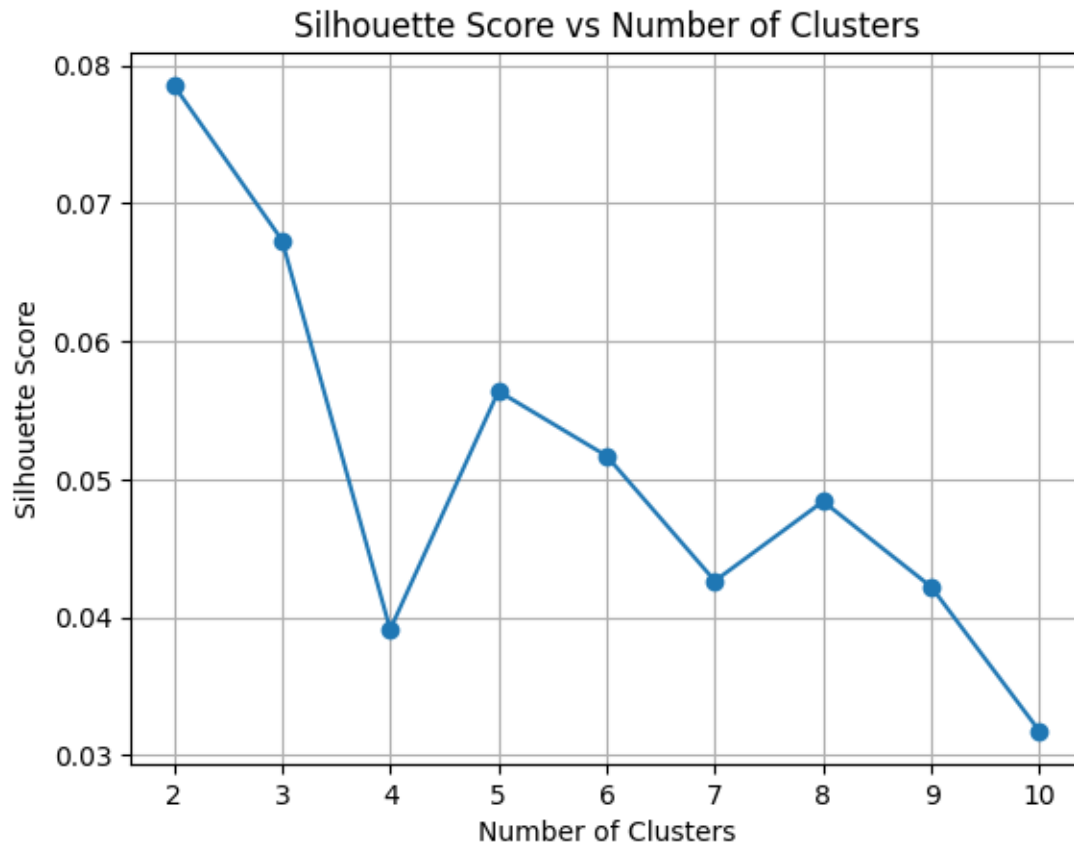
[5 rows x 99 columns]

```
[98]: # Initialize lists to store silhouette scores and cluster numbers
silhouette_scores = []
cluster_numbers = range(2, 11) # Range of cluster numbers to try

# Iterate through each cluster number
for n_clusters in cluster_numbers:
    # Fit KMeans clustering model
    kmeans = KMeans(n_clusters=n_clusters, random_state=42)
    cluster_labels = kmeans.fit_predict(als_df_cleaned)

    # Calculate silhouette score
    silhouette_avg = silhouette_score(als_df_cleaned, cluster_labels)
    silhouette_scores.append(silhouette_avg)

# Plot silhouette scores versus number of clusters
plt.plot(cluster_numbers, silhouette_scores, marker='o')
plt.xlabel('Number of Clusters')
plt.ylabel('Silhouette Score')
plt.title('Silhouette Score vs Number of Clusters')
plt.xticks(cluster_numbers)
plt.grid(True)
plt.show()
```



```
[93]: # Select the optimal number of clusters based on the plot
      optimal_num_clusters = silhouette_scores.index(max(silhouette_scores)) + 2
      print("Optimal number of clusters:", optimal_num_clusters)
```

Optimal number of clusters: 2

The cluster silhouette score is a metric used to evaluate the quality of clusters in K-means clustering. It provides a measure of how similar an object is to its own cluster (cohesion) compared to other clusters (separation). The silhouette score ranges from -1 to 1, where:

A score close to +1 indicates that the object is well matched to its own cluster and poorly matched to neighboring clusters. A score around 0 indicates that the object is on or very close to the decision boundary between two neighboring clusters. A score close to -1 indicates that the object is likely to be assigned to the wrong cluster.

Based on the above plot 2 is the optimal number of cluster as that have the highest silhouette\_scores value.

```
[100]: # Optimal number of clusters chosen from silhouette score
      optimal_num_clusters = 2

      # Fit K-means model with optimal number of clusters
```

```

kmeans = KMeans(n_clusters=optimal_num_clusters, random_state=42)
cluster_labels = kmeans.fit_predict(als_df_cleaned)

# Add cluster labels to the original data
data['Cluster'] = cluster_labels

```

Loaded the dataset and preprocess it by scaling the data. Calculated the silhouette scores for different numbers of clusters and plot them against the number of clusters. Identified the optimal number of clusters based on the plot. Using K-means model calculate the optimal number of clusters.

```

[101]: # Fit a PCA transformation with two features to the scaled data
pca = PCA(n_components=2)
pca_features = pca.fit_transform(scaled_data)

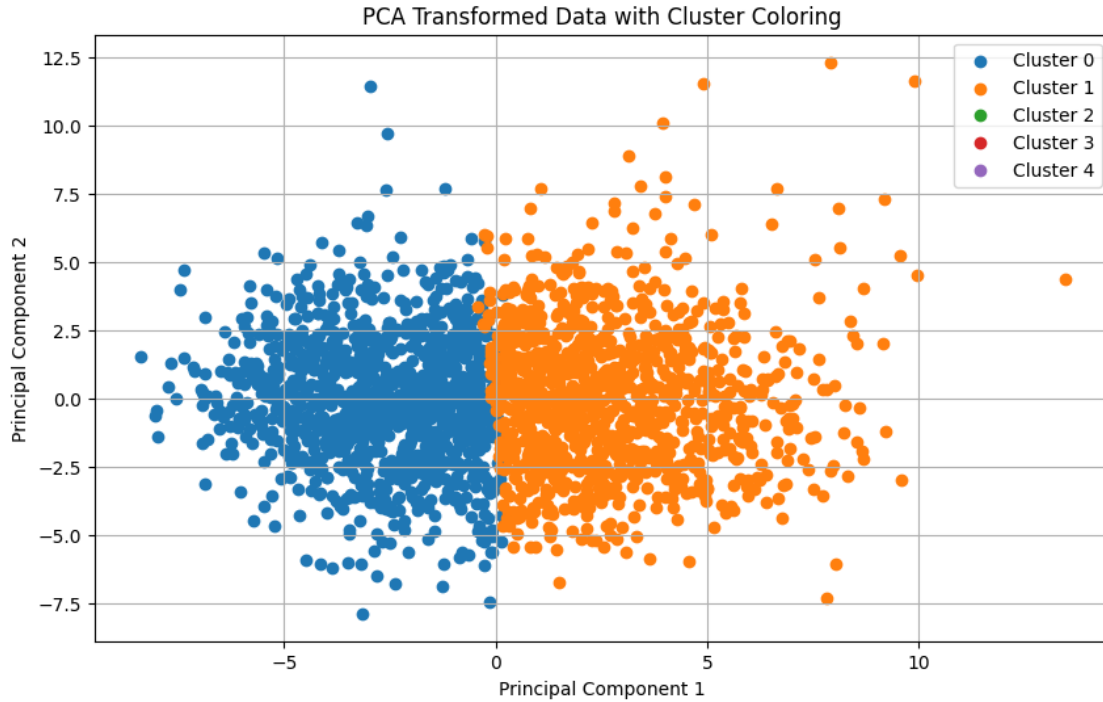
[102]: # Create a DataFrame with PCA features and cluster labels
pca_df = pd.DataFrame(pca_features, columns=['PC1', 'PC2'])
pca_df['Cluster'] = cluster_labels

# Plot scatter plot with cluster coloring
plt.figure(figsize=(10, 6))
for cluster in range(5): # Adjust based on the number of clusters
    cluster_data = pca_df[pca_df['Cluster'] == cluster]
    plt.scatter(cluster_data['PC1'], cluster_data['PC2'], label=f'Cluster_{cluster}')

plt.title('PCA Transformed Data with Cluster Coloring')
plt.xlabel('Principal Component 1')
plt.ylabel('Principal Component 2')
plt.legend()
plt.grid(True)
plt.show()

```





Above code performs all the specified tasks, including data preprocessing, determining the optimal number of clusters, fitting a K-means model, performing PCA transformation, and creating a scatterplot of the PCA-transformed data with cluster coloring.

The scatterplot visualizes the PCA-transformed data, where each point represents a sample from the dataset. The points are colored according to their assigned cluster labels obtained from the K-means clustering algorithm. By examining this plot, we can observe how well the clusters are separated in the reduced feature space (PCA components 1 and 2). If the clusters are well-separated, it indicates that the K-means algorithm has successfully grouped similar samples together. On the other hand, if the clusters overlap significantly, it suggests that the algorithm may not have effectively captured the underlying structure of the data. This visualization provides insights into the distribution and separation of clusters, aiding in the interpretation and evaluation of the clustering results.